

Data Mining Techniques

Databeestjes

Jaro Camphuijsen (jcn690), Rahiel Kasim (rkm700), Jan Westerdiep (jwp800)

April 17, 2016

1 Exploring a small data set

For this task, we decided to use Python with sklearn, as we are proficient in the language and heard good things about the package. The data set is quite nice; we have little missing data values, and most of the attributes are easily parsed. In total, we have 129 data points with 16 attributes, distributed over 13 questions; see Table 1.

What immediately jumps out is the fact that each binary question has its own representation, so we will have to parse this. Moreover, we see that only two people stood up; we will refrain from using attribute 10. Lastly, a few of the enumerative questions had nonlegal answers which we will have to filter out. This results in a few missing data points.

In a recent lecture, [1] a very nice broad exploratory analysis was made; instead of repeating these results, we chose to look at some specific interesting points.

Col	Question	Type
1	Study programme	enum
2-5	Courses (4×)	bool
6	Gender	enum
7	Chocolate makes you...	enum
8	Birthday	date
9	Neighbour count	int
10	Stand up?	bool
11	Stress level	float
12	Money question	float
13	Random number	float
14	Bedtime	time
15-16	Good day?	string

Table 1: The different attributes for the given data set.

1.1 Study programmes

It seems there are only a few different study programmes, or combinations of those, in which students are enrolled. See the left of Figure 1. The majority of the students is enrolled in the programmes Artificial Intelligence, Business Analytics, and Computational Science; this is explained by the fact that Data Mining is an elective course for each of these programmes.

1.2 Birthday problem

Consider the following: in a group of n people, what is the probability that two people share the same birthday? The exact answer [2] to this is

$$p(n) = 1 - \frac{365!}{365^n(365 - n)!}.$$

In our case, this probability amounts to $\approx 1 - 6 \cdot 10^{-12}$; looking at the data set, we see that indeed there are birthday collisions. In fact, we have four collisions on February 14! The probability of this is much harder to compute explicitly [3]; for $n \rightarrow \infty$, a good approximation to this turns out to be

$$p(n, m) = 1 - \exp\left(-\binom{n}{m}/365^{m-1}\right) \implies p(128, 4) \approx 19.7\%,$$

which makes the quadruple occurrence of February 14 much more likely than we had anticipated.

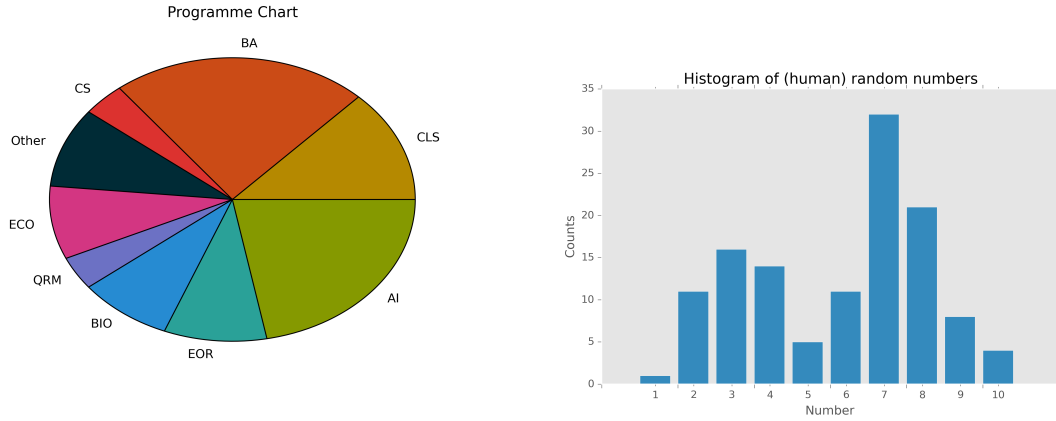


Fig. 1: Left: The distribution of different study programmes of the students enrolled in this course. Right: The distribution of random numbers generated by the students enrolled in this course.

1.3 Random number distribution

Looking at the results of the random number generator, we see that the distribution is not uniform. It is known that humans are generally terrible at generating random numbers, a fact clearly shown at the right of Figure 1. There have been numerous studies on ‘random’ number choice in human psychology, especially on the preference for the number seven. One explanation is from a cognitive and mathematical point of view. When asked to choose a random number between 0 and 10, people tend to maximize the ‘randomness’ of the chosen number; to do this, they look at certain special properties (like being an endpoint 1 or 10, an even number, a power of two) and choose the number with the least of those properties as suggested by Kubova & Psotka (1976). [4]

Another explanation is of a cultural nature, since seven has a special role in the Bible and the Gregorian calendar. This explanation also accounts for the popularity of the number three due to its special significance in the Bible.

The number of occurrences of the number seven over the whole data set is 26%, which is very close to the value of 28.4% found by Kubova & Psotka.

1.4 Basic Classification

We had multiple ideas regarding the classification task we wanted to perform. One of them was to predict the programme a person was taking¹ given the courses they’ve done; the reasoning behind this was that different programmes teach different classes. Our hope that this metric has enough discriminatory power was in vain; classification using a decision tree systematically guessed only one of three programmes.

Other approaches showed similar results, so we abandoned this idea in favour of predicting whether a student had done a specific course given the other three courses. We partitioned the data into training and test data, used a number of different classifiers with default parameters to get a grasp of which ones work well², and then used cross-validation to select optimal parameters on the best two.

Our ‘forest’ of initial classifiers consisted of: a decision tree, a random forest, a Naive Bayes, a k -neighbours classifier, and two support vector machines (one with a linear kernel, and one with a RBF kernel). Taking the target course to be Machine Learning, we toyed with taking different features. Ultimately, the best results were found using just the remaining courses (IR, DB, stats); adding, e.g. gender, stress, or study programme only added confusion.

¹ In our data preprocessing step, we reduced the total number of programmes to nine, setting programmes with less than four students to “Other”, removing students with missing values, and duplicating students with double programmes.

² We used the default ‘mean accuracy’ scoring mechanism contained in sklearn.

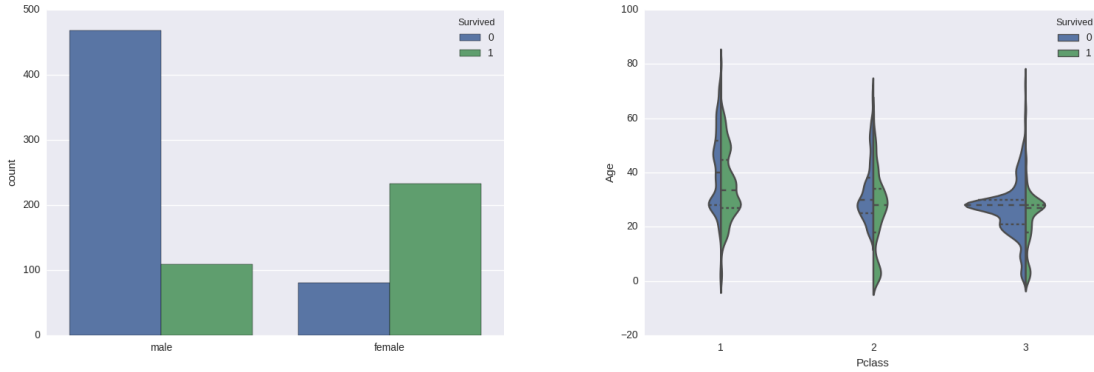


Fig. 2: Left: Number of Titanic survivors depending on sex. Right: Violin graph for the fraction of survivors with age on the y axis and plotted for the three different passenger classes

From these preliminary results, we determined that the Decision Tree and SVM with RBF kernel yielded the best results (of around 72%). We then included some parameter estimation using cross-validation to test a parameter setting, and a grid search to select such settings. The data set however ended up being too small, or our problem too easy, or a combination of both: the best parameters were found to differ from very small to very large in each change of training and data set, with very little change in actual score.

Of course, estimating stress levels would have been more interesting. However, stress values are floating points, hence estimating them is not a classification but a regression task. We decided to pour our remaining time in optimizing the results of task 2 rather than this one.

2 Predicting Titanic Survivors

The online data mining competition site Kaggle [5] launched a competition about the survivors of the MS Titanic. This competition serves a first start for data mining novices and has an extensive tutorial on basic data mining techniques applied to this problem. The data consists of ± 1800 instances of which half is the training set, with attributes given in Table 2. The second half formed the test set and was given without the “Survived” attribute. While exploring the data set several correlations jump out. As can be seen in the left plot of Figure 2 the sex of a passenger is a very good indicator for the survival chance. In general women and children were given a place in the few rescue boats that were present and therefore they had a higher chance of surviving the disaster. A similar correlation can be seen for age and class in the right plot. The majority of the first class passengers survived and the majority of the third class did not. Also the combination of age and class can be a good indicator since most children of class 2 survived, while for children of class 3 there is no apparent correlation.

Col	Question	Type
1	Passenger ID	int
2	Survived	bool
3	Passenger class	enum
4	Name	string
5	Sex	enum
6	Age	int
7	Siblings and Spouse	int
8	Parents and Children	int
9	Ticket	string
10	Fare	float
11	Cabin	string
12	Embarked?	enum

Table 2: The different attributes for the Titanic data set

2.1 Classification

Looking at the score board and forums of this Kaggle competition, we see a few things: firstly, there are a lot of people with perfect 1.0 accuracy scores, and again a large number of submissions with 88%

accuracy. Looking at the scripts these 88%-submissions have, it is clear they all use the same genetic algorithm counting 48 lines of unreadable (and computer-generated) code.

On the forums it can be read that – most likely – the 1.0 scores were achieved by cheating, i.e., looking up the survival status of every person in the test set on the internet. We also read that a score of 80% can be seen as quite good, whereas above 84% is almost unattainable. Moreover, predicting survival on gender alone reaches 76%; using a simple Random Forest model on a selected set of the features 77%. To score higher than this, more advanced models will have to be used.

In the previous exercise we practiced using different classifiers and models using cross-validation. Therefore, the main point of this Kaggle competition is *feature engineering*: combining, scrapping, adding, and otherwise transforming the available data to find features that are predictive of the final survival.

Firstly we have to address the case of missing data. The train set misses data on the Age, Cabin, and Embark columns; the test set on the Age, Cabin, and Fare columns. The Embark and Fare columns contain only a very small number of missing values (2 and 1, resp.) which we can easily estimate. The Cabin column is very hard to do anything with, so we opted to just drop this. The Age column is the most interesting; Kaggle proposes to fill this with the median value,³ either over the whole set or on a per-sex & class basis. We later estimated this Age using a simple Decision Tree regressor.

The Sex and Embarkation columns are enumerative; we can easily convert these to integers. We will (initially) drop the columns Passenger ID, Name, Ticket, and Cabin, because they contain either useless or hard-to-parse data. As for transforming the data into usable features, the following ideas came up:

- Create a **FamilySize** index, being $\text{Parch} + \text{SibSp} + 1$, and drop Parch and SibSp;
- Bin the Age column;
- Try to find families by using Cabin, Ticket, and Name information, together with FamilySize;
- Extract person titles from their Name.

This last item was frankly not ours; many Kaggle submissions use this feature with good results. The reasoning behind the first idea is that parent-child or sibling-spouse relations really don't differ much; their separation could confuse the classifier. The second item came from the class slides. As for idea three, although a genuinely very interesting problem, we chose not to pursue it, as it was infeasible in this time frame.

To test our models and features, we will take the train data set and divide it into three parts, train on the first two and generate a score on the third. This gives us a crude estimate of the Kaggle score we might expect, and allows us to select between different models. After model selection, we train on the entire data set and predict the survival of the remaining test subjects.

This estimated accuracy is indeed very crude: see Table 3. In column (1), we used the features Age, Pclass, FamilySize, Sex, and Fare. In column (2), we added to this a feature Title. In column (3), we swapped Age for AgeBinned, with bins [5, 15, 30, 45]. The fourth experiment is based on (2), with Age predicted using a Decision Tree Regressor trained on Pclass, Sex, Fare, and Title.

We immediately see a few things. Firstly, the normally pretty good SVM classifier is consistently outperformed by the ensemble methods (which is why we didn't include them in later tests). Secondly, the Kaggle score is consistently much lower than our estimated one. In column (2), we see that even though our AdaBoost algorithm was expected to win, the random forest classifier performed quite a lot better. Thirdly, we see that our second idea of binning the ages has disastrous effects on the performance, even though it was expected to work well. Lastly, we see that estimating the age using a regressor is predicted to work much better, but in actuality performs worse than just using a median value.

Discussion and future work We look back on this Kaggle competition with mixed feelings. We didn't do bad, with a 18/37 position (at 9:50 PM on Sunday). But we expected a little more, more so because we already matched our best score of 78.9% in one of our first submissions, not doing any classifier comparisons, model selections, or adding features. This feeling is amplified by the fact that other submissions with comparable features and models yield accuracies in the 80.5%-range. All in all, we are quite happy with the tool we made and the analyses we did, but the score does not reflect our efforts.

³ Given the outlier distribution, median is possibly a better estimator than mean value.

	(1)	(2)	(3)	(4)
Random Forest	81.1% (77.5%)	80.8% (78.9%)	82.2%	84.5% (77.0%)
Gradient Boosting	80.4%	80.8%	82.5% (76.6%)	n.a.
AdaBoost	78.8%	82.2% (77.0%)	81.8%	82.3%
Decision Tree	75.6%	75.8%	n.a.	n.a.
Linear SVM	78.1%	78.8%	n.a.	n.a.
SVM with RBF kernel	79.8%	80.5%	n.a.	n.a.

Table 3: Estimated and true Kaggle score for a variety of feature-classifier combinations. Each classifier has its parameters tuned by means of a grid search using 5-fold cross-validation; each estimated result was found by training on two-thirds of the available data and testing on the rest.

Had we had more time, we would’ve loved to look more into our third idea above; extracting families from the data is (up to some extent) possible. We can even verify our findings using the Parch and SibSp columns (and possibly the use of the internet).

3 Research and Theory

3.1 State of the art solutions

The “Galaxy Zoo” data mining competition on Kaggle [6] was set up to identify images of galaxies made by the Sloan Digital Sky Survey [7] and Hubbles CANDELS [8] by their shape. Previously the classification of these images was done in a crowd sourcing project where hundreds of volunteers classified galaxies using a search tree of questions. In fact the actual goal of the challenge is to predict the outcome of such a classification by the crowd and therefore it is a regression problem rather than a classification problem.

The competition ended on April 4th and the prize money went to the top three competitors. All three of them used a similar approach with a Convolutional Neural Network (CNN) and similar preprocessing. The Root Mean Square Error (RMSE) of the test set prediction was used as leaderboard score; the winner ended with a RMSE of 0.07467.

The concept of CNN is inspired by the visual cortex of animals, where the neurons all have a receptive field that overlaps with other the receptive fields of other neurons. What a CNN does is to Machine Learn the appropriate filter to extract features from a certain type of data. This in contrast with extracting image features with a certain predefined filter. Because of the nature of the images in this data set, regular features used in computer vision would’ve returned bad results because they are generally accustomed to objects as we find them on earth.

After the maximizing the number of found features, the filter can be applied to the images to extract the features and recognize certain structures of galaxies.

3.2 Measuring Numerical Error

Classification algorithms predict a mapping from a set of attributes to a corresponding discrete value or label. It is easy to evaluate their performance: we simply check whether the predicted label was correct or not. A prediction is simply false or true. Regression algorithms predict numerical or continuous values: there is a continuous spectrum of errors, so we need a different measure of performance. Here we will discuss two such error measures: mean squared error (MSE) and mean absolute error (MAE). If p_1, \dots, p_n are the predicted values and a_1, \dots, a_n the actual values, then [9]:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (p_i - a_i)^2 \qquad \text{MAE} = \frac{1}{n} \sum_{i=1}^n |p_i - a_i|$$

The model with the lowest MSE or MAE will then be the best performing model, i.e. the model with the best predictions. The main difference between MSE and MAE is the square in the MSE: the

mean-squared error exaggerates the effect of outliers in the data. MAE does not have this effect, there all error sizes are treated equally. [9] Using MAE is recommended if many outliers are expected when evaluating the numerical performance of regression algorithms. Otherwise there is no difference for model evaluation.

There are also situations where the MSE and MAE are exactly the same. This happens when the terms in the summation are the same: $(p_i - a_i)^2 = |p_i - a_i| \implies p_i = a_i \vee p_i = a_i - 1$ for all i . For example, if $\mathbf{a} = (1, 2, 3)^T$ and $\mathbf{p} = (1, 1, 4)^T$, then $\text{MSE} = \text{MAE} = 2/3$.

3.3 Analysis of a less obvious data set

We are given a set consisting of English text messages, each text labeled as either ‘spam’ or ‘ham’, with the latter signifying a bona fide SMS.

Text mining is a very active area of both research and business. Our specific task within text mining is akin to *sentiment estimation*, and multiple Kaggle competitions have been done with this exact task in mind. Most submissions seem to use a so-called *bag-of-words* approach; instead of trying to parse the natural language contained in each text message, we create a ‘vocabulary’ or dictionary of every word contained in this set of texts.⁴ The resulting histograms (defined as the number of each word a text contains) then become our feature vectors. As we did before, we subdivide the data set into a training and test set, retaining one third of the total size for testing.⁵

In practice, using these histograms as-is yields some inconsistencies: consider a text message S , with some distribution D . Then the message $SSSS$ (four repetitions of S) has distribution $4D$, which is very different from D , although the two messages discuss the same topic and will hence carry the same label. To counter this, well-performant Kaggle submissions will often divide these histograms by their sum, yielding something like a “frequency histogram”.

A large variety of classifiers can be used on tasks like these. Looking at the sklearn flow chart, [10] we see that the SVM with a linear kernel, Stochastic GD classifier, and Naive Bayes are good contenders. Again, these all have their respective parameters which we can optimize using a grid search and cross-validation.

Running the above method in its most simple form yields 8713 features per sample, with results visible in the column (1) of Table 4. The other columns show results after distillation of the features: (2) we have 4213 features left if we remove all words that occur in only one text message; (3) 3965 features result after removing *stop words* (e.g., ‘a’, ‘the’, ‘which’) from the texts before tokenizing.

Looking at this table, we see that removing single-document words makes hardly any difference in terms of accuracy, and that removing the stop words doesn’t do much either. As far as model selection goes, we see that the linear SVM and Stochastic Gradient Descent classifiers perform best, with comparable results. This is again in accordance with the previously mentioned flow chart.

If we think about how the linear SVM works, we can infer some interesting points from the weights. [11] Large absolute values have a big say in the final class prediction; negative weights correspond to the negative class (in our case ham) and vice versa. Using the linear SVM-model that yielded the best results, we can come up with two lists of the ‘hammiest’ and ‘spammiest words’, which are pretty much in line with what one would expect.

	(1)	(2)	(3)
# features	8713	4213	3965
DecisionTree	96.5%	96.3%	96.1%
Linear SVM	98.2%	98.4%	98.0%
Stochastic GD	98.2%	98.4%	98.2%
Multinomial NB	97.2%	98.1%	98.2%

Table 4: Accuracy on a held-out test set for different classifiers, for different feature settings. All classifiers have their parameters optimized by cross-validation, and each accuracy is found by averaging over 10 experiments.

⁴ A few things are important here: firstly, the definition of a *word*. Most often (as is in our case), a sequence of two or more alphanumeric characters surrounded by non-alphanumeric characters is deemed a word. Moreover, these words need to be tokenized, which means given an index. Often times people choose to convert every word to lowercase, but other transformations are also possible.

⁵ In order to reduce noise, we repeat each experiment 10 times, and average the result.

Hammiest words liked, happy, mail, yup, ok, morning,
cool, cancer

Spammiest words uk, www, claim, service, filthy, win, games

References

- [1] M. Hoogendoorn. Data mining techniques; lecture 2. University Lecture, 2016.
- [2] Wikipedia. Birthday problem — Wikipedia, The Free Encyclopedia, 2016. [Online; accessed 11-April-2016].
- [3] Anirban DasGupta. The matching, birthday and the strong birthday problem: a contemporary review, 2003.
- [4] Kubovy M. and Psotka J. The predominance of seven and the apparent spontaneity of numerical choices. *Journal of Experimental Psychology: Human Perception and Performance*, 2(2):291–294, 1976.
- [5] Kaggle inc. Titanic: Machine learning from disaster, 2016. [<https://www.kaggle.com/c/titanic>; accessed 15-April-2016].
- [6] Sander Dieleman. My solution for the galaxy zoo challenge, 2016. [<http://benanne.github.io/2014/04/05/galaxy-zoo.html>; accessed 16-April-2016].
- [7] Astrophysical Research Consortium. Sloan Digital Sky Survey. [<http://www.sdss.org/>; accessed 16-April-2016].
- [8] Astrophysical Research Consortium. Cosmic Assembly Near-infrared Deep Extra-galactic Legacy Survey. [<http://candels.ucolick.org/>; accessed 16-April-2016].
- [9] Mark A. Hall Ian H. Witten, Frank Eibe. *Data mining : practical machine learning tools and techniques*. Elsevier, third edition, 2011.
- [10] Scikit-learn. Choosing the right estimator. [http://scikit-learn.org/stable/tutorial/machine_learning_map/index.html; accessed 16-April-2016].
- [11] Stack Exchange. How does one interpret SVM feature weights? [<http://stats.stackexchange.com/questions/39243>; accessed 16-April-2016].